

```
In [5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression, LinearRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix, classification_report

from sklearn.model_selection import train_test_split

import warnings
warnings.filterwarnings('ignore')
```

```
In [6]: ds=pd.read_csv('salary_project.csv')
```

```
In [9]: ds.head()
```

```
Out[9]:
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
0	Prof	B	19	18	Male	139750
1	Prof	B	20	16	Male	173200
2	AsstProf	B	4	3	Male	79750
3	Prof	B	45	39	Male	115000
4	Prof	B	40	41	Male	141500

```
In [11]: ds.shape
```

```
Out[11]: (397, 6)
```

```
In [13]: ds.dtypes
```

```
Out[13]: rank                object
discipline                object
yrs.since.phd             int64
yrs.service               int64
sex                       object
salary                    int64
dtype: object
```

```
In [14]: ds.columns
```

```
Out[14]: Index(['rank', 'discipline', 'yrs.since.phd', 'yrs.service', 'sex', 'salary'],
dtype='object')
```

```
In [16]: ds.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 397 entries, 0 to 396
Data columns (total 6 columns):
#   Column                Non-Null Count  Dtype
---  -
0   rank                  397 non-null   object
```

```

1  discipline      397 non-null    object
2  yrs.since.phd   397 non-null    int64
3  yrs.service     397 non-null    int64
4  sex             397 non-null    object
5  salary          397 non-null    int64
dtypes: int64(3), object(3)
memory usage: 18.7+ KB

```

In [17]:

```
ds.isnull().sum()
```

Out[17]:

```

rank          0
discipline    0
yrs.since.phd 0
yrs.service   0
sex           0
salary        0
dtype: int64

```

In [18]:

```
ds.describe()
```

Out[18]:

	yrs.since.phd	yrs.service	salary
count	397.000000	397.000000	397.000000
mean	22.314861	17.614610	113706.458438
std	12.887003	13.006024	30289.038695
min	1.000000	0.000000	57800.000000
25%	12.000000	7.000000	91000.000000
50%	21.000000	16.000000	107300.000000
75%	32.000000	27.000000	134185.000000
max	56.000000	60.000000	231545.000000

In [37]:

```
ds.loc[ds['salary']==' ']
```

Out[37]:


```
rank  discipline  yrs.since.phd  yrs.service  sex  salary
```

In [19]:

```
sns.countplot(x='sex', data=ds)
```

Out[19]:

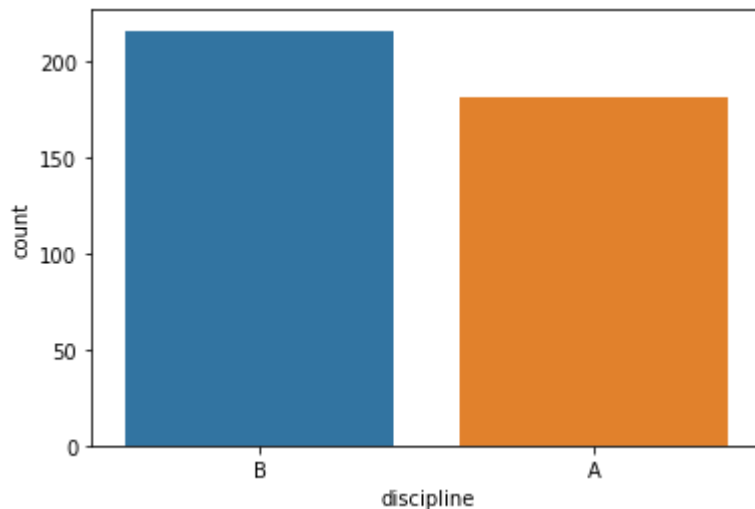
```
<AxesSubplot:xlabel='sex', ylabel='count'>
```



The gender in the data is not evenly distributed

```
In [25]: sns.countplot(x='discipline',data=ds)
```


```
Out[25]: <AxesSubplot:xlabel='discipline', ylabel='count'>
```



There is a fair distribution in discipline

```
In [26]: sns.countplot(x='rank',data=ds)
```


```
Out[26]: <AxesSubplot:xlabel='rank', ylabel='count'>
```



The data has is heavily skewed in favor of professors as against assistant and associate professors

```
In [22]: sns.distplot(ds['salary'])
```


```
Out[22]: <AxesSubplot:xlabel='salary', ylabel='Density'>
```



The data does not follow a fair distribution curve

```
In [23]: sns.distplot(ds['yrs.service'])
```


```
Out[23]: <AxesSubplot:xlabel='yrs.service', ylabel='Density'>
```



The data is not fairly distrinuted

```
In [24]: sns.distplot(ds['yrs.since.phd'])
```

```
Out[24]: <AxesSubplot:xlabel='yrs.since.phd', ylabel='Density'>
```



The data seem well distributed

```
In [30]: ds=pd.DataFrame(ds)
ds
```

```
Out[30]:
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
0	Prof	B	19	18	Male	139750
1	Prof	B	20	16	Male	173200
2	AsstProf	B	4	3	Male	79750
3	Prof	B	45	39	Male	115000
4	Prof	B	40	41	Male	141500
...
392	Prof	A	33	30	Male	103106
393	Prof	A	31	19	Male	150564
394	Prof	A	42	25	Male	101738
395	Prof	A	25	15	Male	95329
396	AsstProf	A	8	4	Male	81035

397 rows × 6 columns

```
In [38]: from sklearn.preprocessing import OrdinalEncoder
enc=OrdinalEncoder()
```

```
In [39]: for i in ds.columns:
          if ds[i].dtypes=='object':
              ds[i]=enc.fit_transform(ds[i].values.reshape(-1,1))
```

```
In [40]: ds
```

```
Out[40]:
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
0	2.0	1.0	19	18	1.0	139750
1	2.0	1.0	20	16	1.0	173200

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
2	1.0	1.0	4	3	1.0	79750
3	2.0	1.0	45	39	1.0	115000
4	2.0	1.0	40	41	1.0	141500
...
392	2.0	0.0	33	30	1.0	103106
393	2.0	0.0	31	19	1.0	150564
394	2.0	0.0	42	25	1.0	101738
395	2.0	0.0	25	15	1.0	95329
396	1.0	0.0	8	4	1.0	81035

397 rows × 6 columns

In [42]:

```
ds.dtypes
```

Out[42]:

```
rank          float64
discipline    float64
yrs.since.phd  int64
yrs.service    int64
sex           float64
salary        int64
dtype: object
```

In [45]:

```
ds.describe()
```

Out[45]:

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
count	397.000000	397.000000	397.000000	397.000000	397.000000	397.000000
mean	1.508816	0.544081	22.314861	17.614610	0.901763	113706.458438
std	0.757486	0.498682	12.887003	13.006024	0.298010	30289.038695
min	0.000000	0.000000	1.000000	0.000000	0.000000	57800.000000
25%	1.000000	0.000000	12.000000	7.000000	1.000000	91000.000000
50%	2.000000	1.000000	21.000000	16.000000	1.000000	107300.000000
75%	2.000000	1.000000	32.000000	27.000000	1.000000	134185.000000
max	2.000000	1.000000	56.000000	60.000000	1.000000	231545.000000

In [41]:

```
ds.corr()
```

Out[41]:


	rank	discipline	yrs.since.phd	yrs.service	sex	salary
rank	1.000000	-0.086266	0.525500	0.447499	0.132492	0.522207
discipline	-0.086266	1.000000	-0.218087	-0.164599	0.003724	0.156084
yrs.since.phd	0.525500	-0.218087	1.000000	0.909649	0.148788	0.419231
yrs.service	0.447499	-0.164599	0.909649	1.000000	0.153740	0.334745
sex	0.132492	0.003724	0.148788	0.153740	1.000000	0.138610

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
salary	0.522207	0.156084	0.419231	0.334745	0.138610	1.000000

In [54]:

```
plt.figure(figsize=(22,7))
sns.heatmap(ds.describe()[1:],annot=True,linewidths=0.1,linecolor='white',fmt='%.2f')
```


Out[54]: <AxesSubplot:>



In [52]:

```
plt.figure(figsize=(22,7))
sns.heatmap(ds.corr(),annot=True,linewidths=0.1,linecolor='white',fmt='0.2f')
```

Out[52]: <AxesSubplot:>



From the above, salary seems not to have a good correlation with the other variables. The correlation is seen to be lowest with sex and discipline.

In [55]:

```
# Correlation with the target column:
```

In [56]:

```
ds.corr()['salary'].sort_values()
```

Out[56]:

```
sex          0.138610
discipline   0.156084
yrs.service  0.334745
yrs.since.phd 0.419231
rank         0.522207
salary       1.000000
Name: salary, dtype: float64
```

In [58]:

```
ds.skew().sort_values(ascending=False)
```

```
Out[58]: salary          0.714568  
yrs.service      0.650569  
yrs.since.phd    0.300880  
discipline      -0.177684  
rank            -1.151164  
sex             -2.709958  
dtype: float64
```


Keeping +/-0.65 as the range for skewness, here are the columns which does not lie within this range;

- salary
- rank
- sex

```
In [59]: #Checking Outliers:
```


```
In [60]: ds['salary'].plot.box()
```

```
Out[60]: <AxesSubplot:>
```




```
In [61]: ds['yrs.service'].plot.box()
```

```
Out[61]: <AxesSubplot:>
```



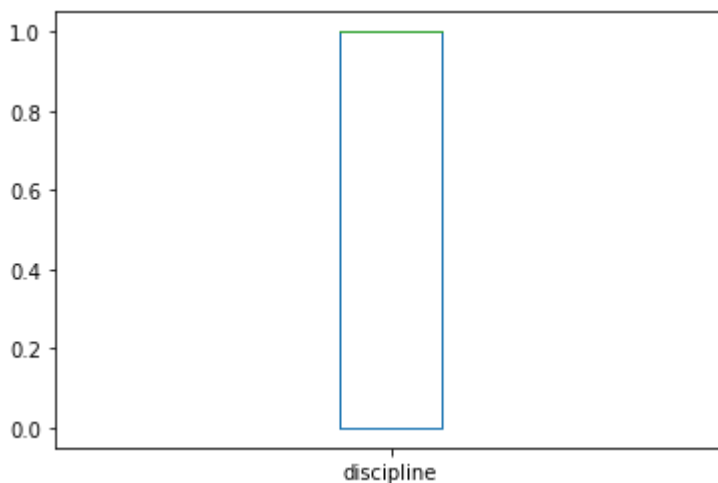

```
In [62]: ds['yrs.since.phd'].plot.box()
```

```
Out[62]: <AxesSubplot:>
```




```
In [63]: ds['discipline'].plot.box()
```

```
Out[63]: <AxesSubplot:>
```




```
In [64]: ds['rank'].plot.box()
```

```
Out[64]: <AxesSubplot:>
```



```
In [65]: ds['sex'].plot.box()
```

Out[65]: <AxesSubplot:>



In [66]: `ds.shape`

Out[66]: (397, 6)

In [67]: `#Removing Outliers:`

In [69]: `from scipy.stats import zscore
import numpy as np
z=np.abs(zscore(ds))
threshold=3
np.where(z>3)`

Out[69]: (array([9, 19, 24, 34, 35, 43, 47, 48, 52, 63, 68, 84, 90,
103, 114, 119, 123, 127, 131, 132, 133, 148, 153, 179, 186, 218,
230, 231, 233, 237, 245, 253, 254, 274, 316, 323, 330, 332, 334,
341, 358, 361, 364]),
array([4, 4, 4, 4, 4, 5, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 4, 4,
4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 3, 4, 4, 4, 4, 4, 5]))

In [70]: `ds_new=ds[(z<3).all(axis=1)]`

In [71]: `ds.shape`

Out[71]: (397, 6)

In [72]: `ds_new.shape`

Out[72]: (354, 6)

In [73]: `x=ds_new.iloc[:,0:-1]`

In [74]: `x`

Out[74]: `rank discipline yrs.since.phd yrs.service sex`

	rank	discipline	yrs.since.phd	yrs.service	sex
0	2.0	1.0	19	18	1.0
1	2.0	1.0	20	16	1.0
2	1.0	1.0	4	3	1.0
3	2.0	1.0	45	39	1.0
4	2.0	1.0	40	41	1.0
...
392	2.0	0.0	33	30	1.0
393	2.0	0.0	31	19	1.0
394	2.0	0.0	42	25	1.0
395	2.0	0.0	25	15	1.0
396	1.0	0.0	8	4	1.0

354 rows × 5 columns

In [75]: `y=ds_new.iloc[:, -1]`

In [76]: `y`

Out[76]:

```

0      139750
1      173200
2       79750
3     115000
4     141500
...
392    103106
393    150564
394    101738
395     95329
396     81035
Name: salary, Length: 354, dtype: int64
```

In [77]: `x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=20,random_state=`

In [78]: `x_train.shape`

Out[78]: (334, 5)

In [79]: `x_test.shape`

Out[79]: (20, 5)

In [80]: `y_train.shape`

Out[80]: (334,)

In [81]: `y_test.shape`

Out[81]: (20,)

In [82]: `lr=LinearRegression()`

In [83]: `lr.fit(x_train,y_train)`

Out[83]: `LinearRegression()`

In [84]: `ds.columns`

Out[84]: `Index(['rank', 'discipline', 'yrs.since.phd', 'yrs.service', 'sex', 'salary'], dtype='object')`

In [85]: `pred=lr.predict(x_test)`
`print('Predicted Salary: ',pred)`
`print('Actual Salary: ',y_test)`

```
Predicted Salary: [119800.41922748 117784.53516541 122164.23129701  89609.339
37947
 117961.82571662 126641.2874388  126111.94560093 115680.00582773
 100657.20462655 124007.41626325 131552.85522404 105566.24259602
 113659.53031027 104338.98310365 118839.09556194 117695.8898898
 122688.9816795  103788.60453441 123130.14641794  88361.5113319 ]
Actual Salary:  22      93904
305      111350
85       132825
376       74856
113      104279
238       77202
320      104428
116      148500
61       75243
155      118971
4        141500
11       79800
110      112429
83       88825
250      109000
105      113543
235       81700
162       98510
44       94384
54       103760
Name: salary, dtype: int64
```

In []: